

Ribhav Kapur

647-740-8459

ribhav.kapur99@gmail.com

Annual GPA: 3.79

I am a 4th year Computer Science student focusing in Artificial Intelligence at the University of Toronto. I am extremely interested in doing research in AI and formulating new learning methods to create smarter and more efficient architectures. My goal is to be able to create an AI being that is conscious (ofcourse the definition of consciousness itself is a very deep question) and that gathers some sort of a deep emotional response from humans. I am extremely fascinated by the brain, consciousness, and how we perceive everything around us and how and why we feel the emotions that we do and how we learn things over our lifetime. I have found that the best way for me to learn something is to build it and the simple complexity of the math behind machine learning really fascinates me.

I believe that when major technological advancement takes place, the benefits of that shift are reaped for years to come as other technologies advance and I believe that is the case with modern AI. I believe that a new shift is required to take AI to the next level and to create artificial consciousness and artificial general intelligence and I would be honoured to be part of the research that makes it happen.

I have worked in AI projects since my first year where I placed top 5 in the Orbis Challenge hackathon where the goal was to develop an algorithm that plays splix.io. I had my first introduction to various search algorithms here like BFS, DFS, A*, DIBFS (Depth-Iterative BFS), etc. The following year, I placed 4th at MakeHarvard for a security drone that leveraged the power of computer vision to autonomously secure areas.

My first research experience was this summer when I wrote a paper about Cycle-GANs under the guidance of Dr. Morteza Rezaeejad. The paper focused on trying to generate a human face from a cartoon. I learnt a lot about stabilising GANs and the importance of preprocessing data from this. My biggest take away from this was that there is a lot more to machine learning than what is taught in class and the only way to truly understand the depths of it is to get first hand experience in machine learning.

Currently, I'm working on a paper on algorithmic trading. I am benchmarking the performance of various machine learning algorithms, both deep learning and non-deep learning, to see how they perform against some strategies written by quantitative analysts. I am using this opportunity to learn about the breadth of machine learning since I feel like these days everyone just gets lost in one particular type of neural network and never looks beyond it. I would like to explore the different methods like decision trees with random forests, XG boosting, bagging SVMs, etc along with regular multi-layer-perceptron networks, larger residual networks, CNNs, RNNs, LSTMs, etc. I am using this as an opportunity to learn more about these different methods and get experience working with them. I am also using a paper trading platform to make trades and see how each model performs (alongside my trading strategy).

I have some experience in Software Development as well. I have coded a security drone from scratch in C# using an arduino as a flight controller for a hackathon at Harvard where I placed top 5. I have also worked in 2 startups, building the tech stacks from scratch. I am proficient in Python, C, JavaScript and I am also comfortable with Java.

